

LDPC decoder on Nvidia GPU

...

Nathan BIETTE • Jean-Raphaël CORNEL • Nicolas LE CLERC • Camille MORIN •
Romain PARRACONE • Samuel PIERRE • Emmanuelle RISSON • Louis ROMAN

1. LDPC decoder algorithm

...

What is a Parity Check code?

- Code words are of fixed length **n**
- In each code word there are **r** check symbols and **n-r** information symbols
- Each symbol in a code word satisfies **r** equations
- Parity Check Matrix

$$n = 7 \quad r = 3$$

$$C = C_1 C_2 C_3 C_4 C_5 C_6 C_7$$

$$\begin{aligned} C_1 \oplus C_2 \oplus C_5 &= 0 \\ C_1 \oplus C_2 \oplus C_4 \oplus C_6 &= 0 \\ C_1 \oplus C_3 \oplus C_4 \oplus C_7 &= 0 \end{aligned}$$

$$M = \begin{bmatrix} 1 & 1 & 0 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 1 & 0 & 0 & 1 \end{bmatrix}$$

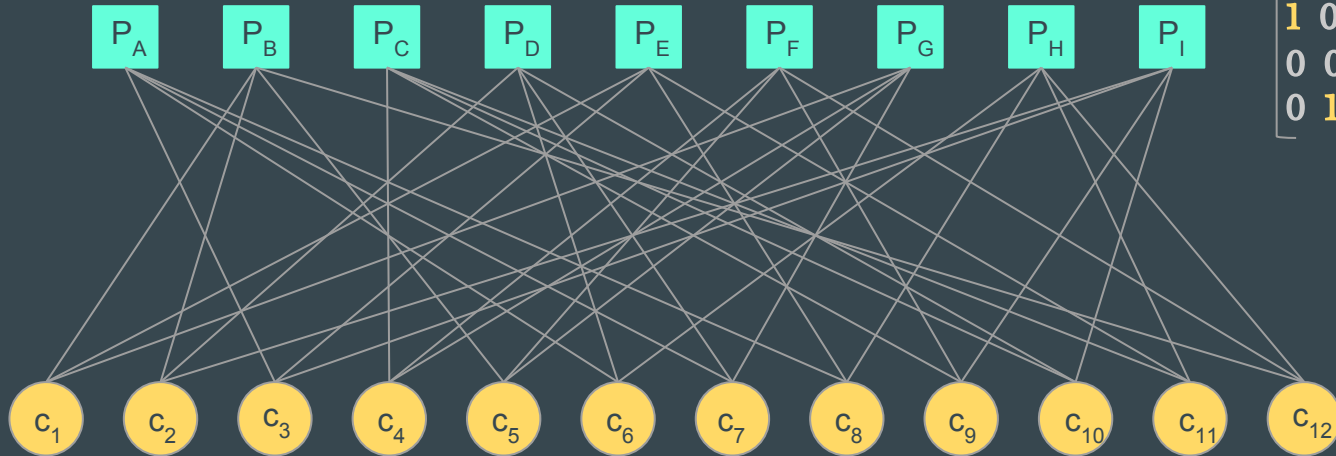
$$M \cdot c^T = 0$$

What is a **Low-Density** Parity Check code?

- Every code digit is in the same number of equations
- Every equation has the same number


$$M = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 & 1 & 1 & 0 \\ 0 & 1 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 \end{bmatrix}$$

Here is a graphical representation:



Estimation of Probabilities by Parity Nodes

For example, 1st equation: $c_3 \oplus c_6 \oplus c_7 \oplus c_8 = 0$



p_i : estimation of $P[c_i=1]$


The parity node:

- receives the estimate p_3, p_6, p_7, p_8
- calculates new estimate $p_{3A}, p_{6A}, p_{7A}, p_{8A}$

$$p_{3A} = p_6 p_7 p_8 + p_6 (1-p_7)(1-p_8) + p_7 (1-p_6)(1-p_8) + p_8 (1-p_6)(1-p_7)$$

Estimation of Probabilities by Bit Nodes

For example, 1st equation: $c_3 \oplus c_6 \oplus c_7 \oplus c_8 = 0$




p_i : estimation of $P[c_i=1]$

The bit node:

- receives three estimate p_{3A}' , p_{3E}' and p_{3I}' from P_A , P_E and P_I
- calculates three new estimate p_{3A} , p_{3E} and p_{3I} for each parity node P_A , P_E and P_I

$$p_{3A} = K p_3 p_{3E} p_{3I} \quad p_{3E} = K p_3 p_{3A} p_{3I} \quad p_{3I} = K p_3 p_{3A} p_{3E}$$

Iterating the process



p_i converge \Rightarrow good estimation of bit node value c_i (hard decision)

To sum up: Low-Density Parity-Check decoder

- Error-correction code
- Near-capacity performance
 - Needs:
 - High throughput
 - Multi-standard support
 - Flexibility and scalability

2. Adjustments of the algorithm

...

Why use a GPU for the LDPC decoder?

- To achieve high throughput
 - GPU (Graphics Processing Unit)
 - high computational throughput
 - several processors working in parallel
 - LDPC
 - massively parallel algorithm
 - clear architecture to divide the operations

Algorithm 2 Depth-first multi-stream scheduling.

```
1: for  $i = 0$  to  $N_{Stream} - 1$  do
2:   memcpyAsync(streams[i], host→device);
3:   for  $j = 0$  to  $N_{iter} - 1$  do
4:     CNP_kernel(streams[i]);
5:     VNP_kernel(streams[i]);
6:   end for
7:   memcpyAsync(streams[i], device→host);
8: end for
9: for  $i = 0$  to  $N_{Stream} - 1$  do
10:  streamSynchronize(streams[i]);
11: end for
```





Fig. 4. Multi-stream LDPC decoding.

Pinned vs. non-pinned memory


- Pinned memory is **more efficient** to achieve data transfer to GPU
- We put the data in a special memory space such that it is directly accessible by the DMA without intervention of the CPU
- If the data is not stored in pinned memory, the program will itself copy data in a pinned memory before sending it to the GPU
- We can directly allocate the arrays in pinned host memory to **reduce time** of the copy

Use: `cudaMallocHost()`

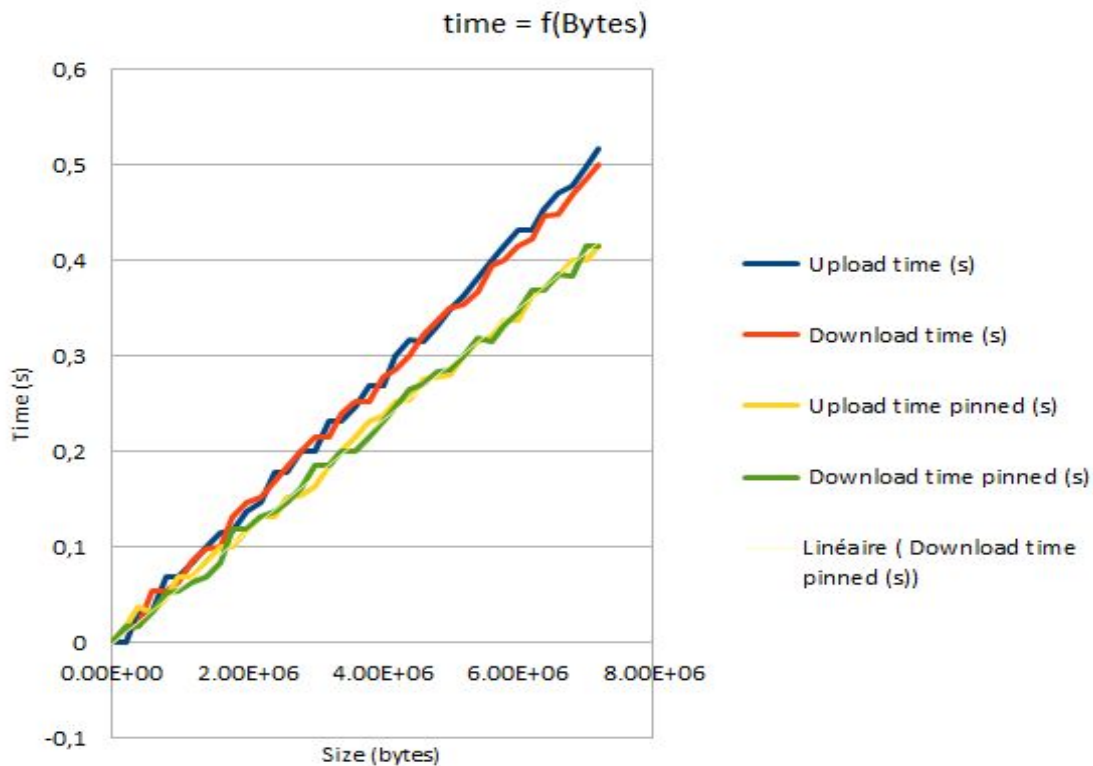
Pageable Data Transfer



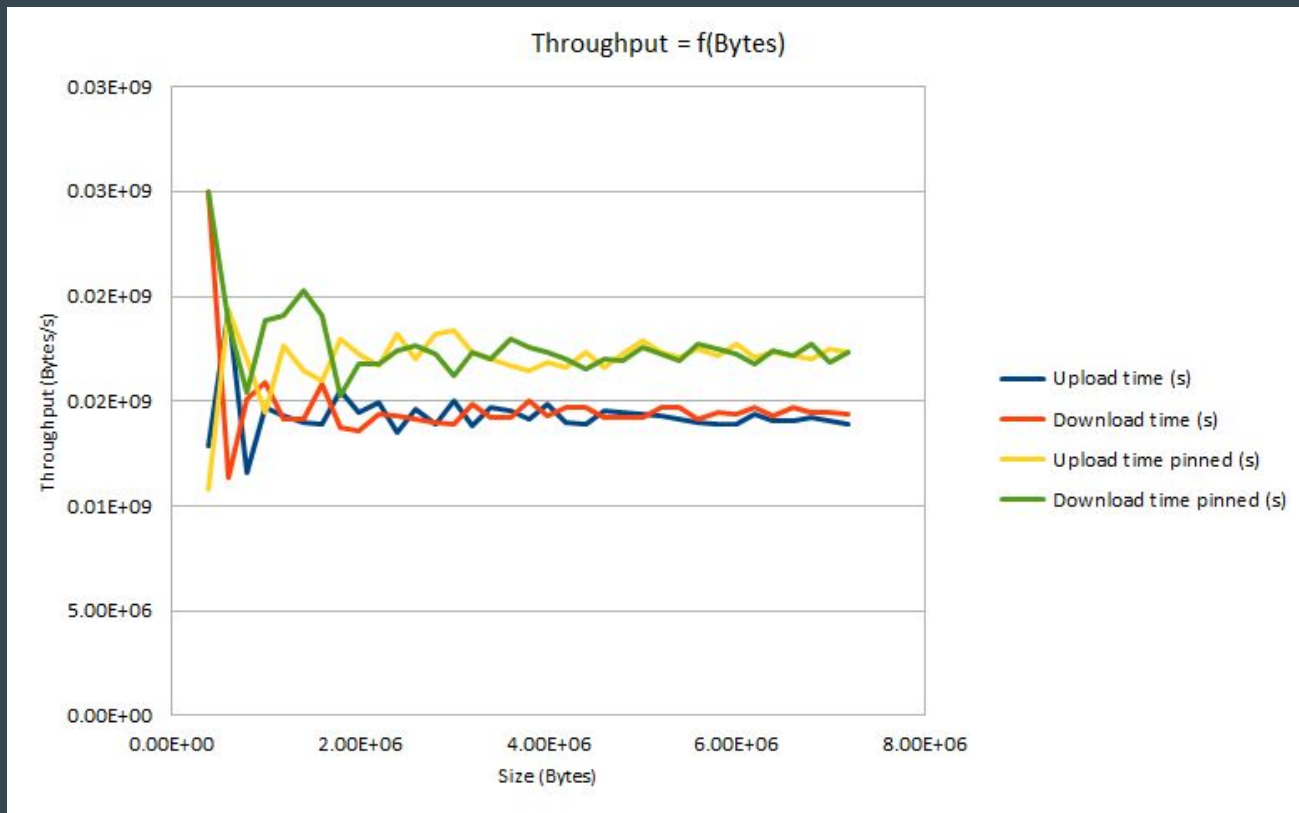
Pinned Data Transfer



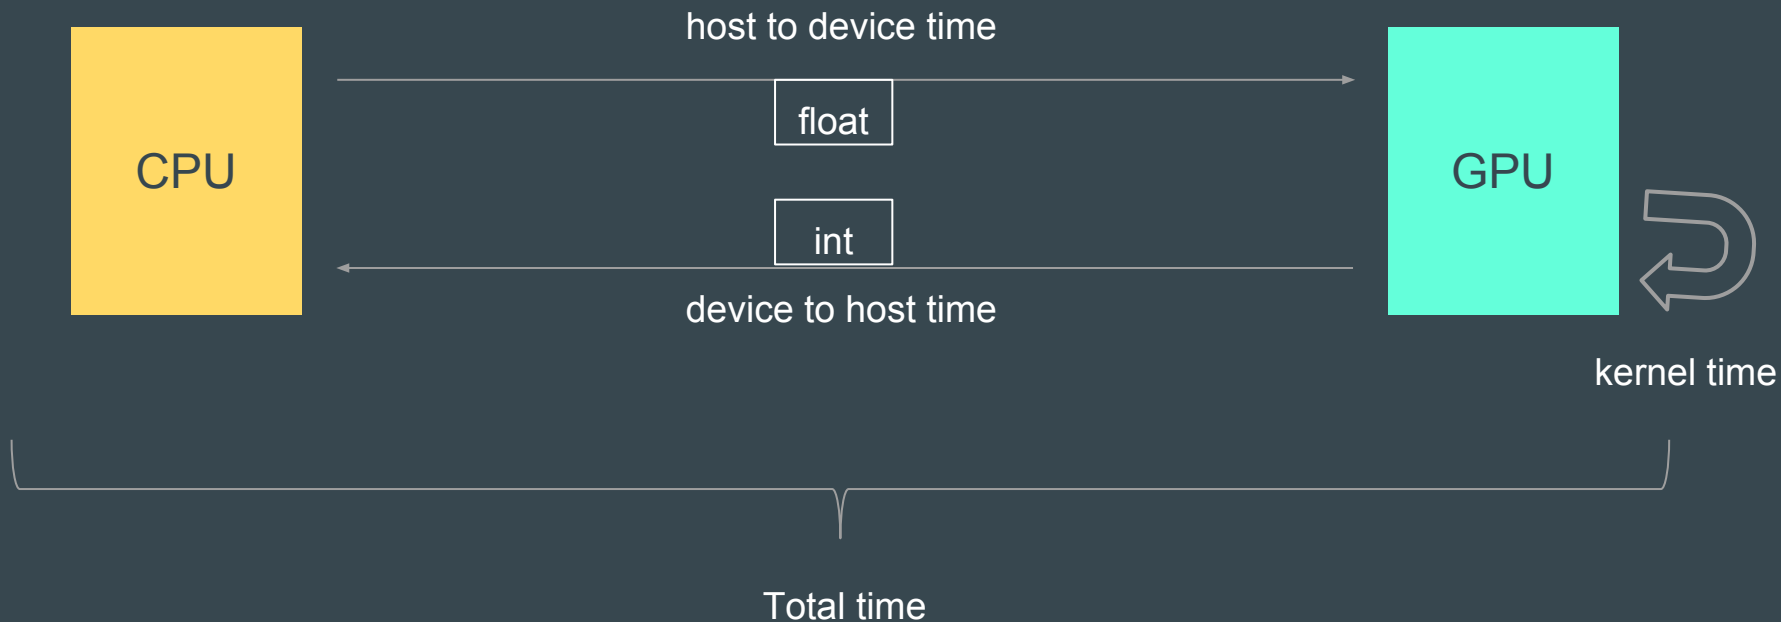
Comparison between the two memories



Comparison between the two memories



Timing measurements




3. Benchmark results

...

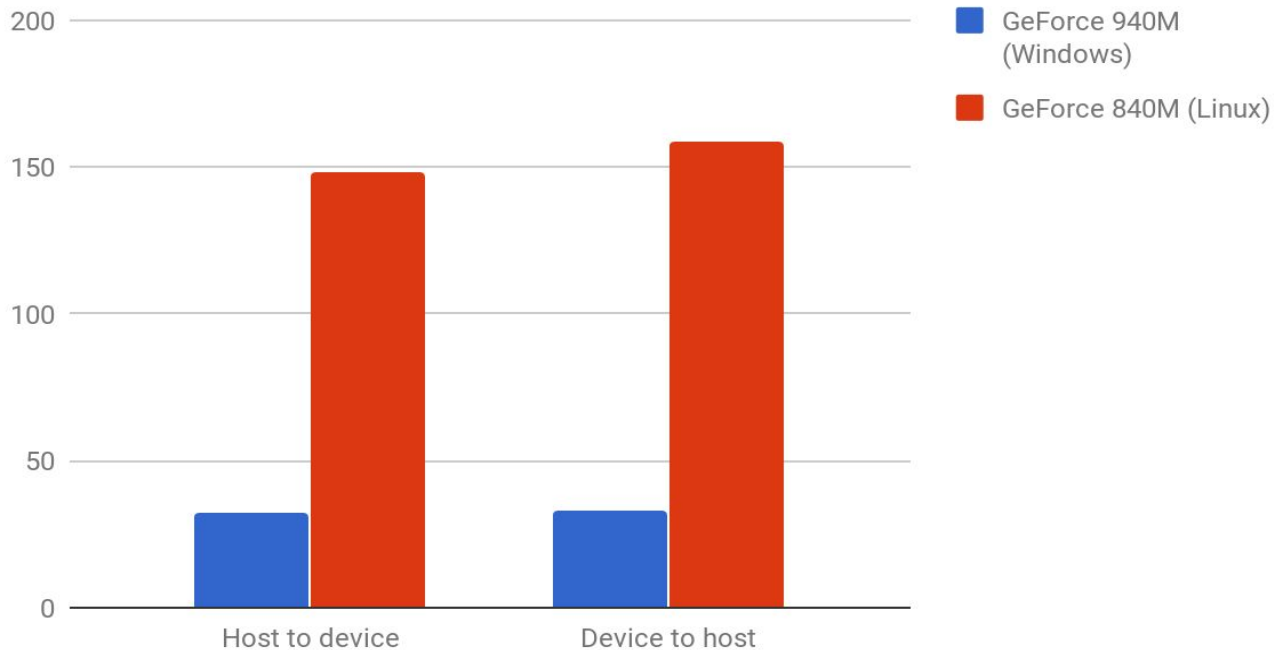
Measures

Throughput of the LDPC algorithm in Mb/s

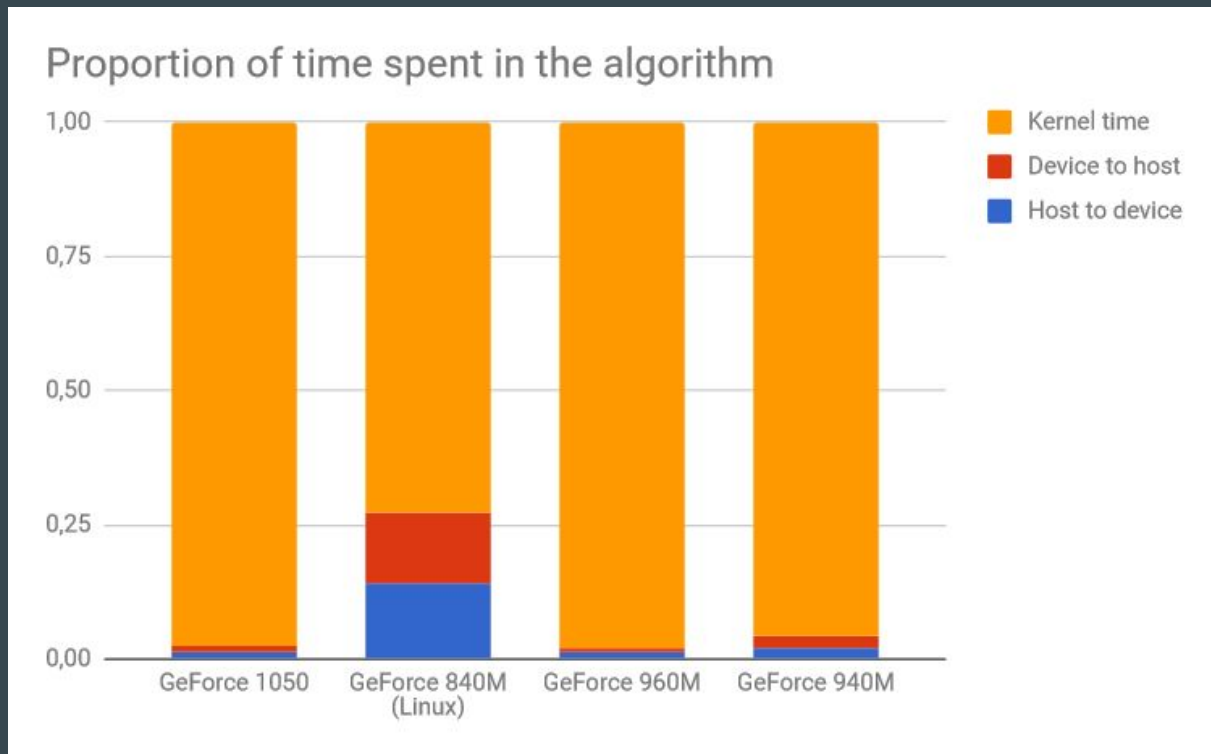


Measures

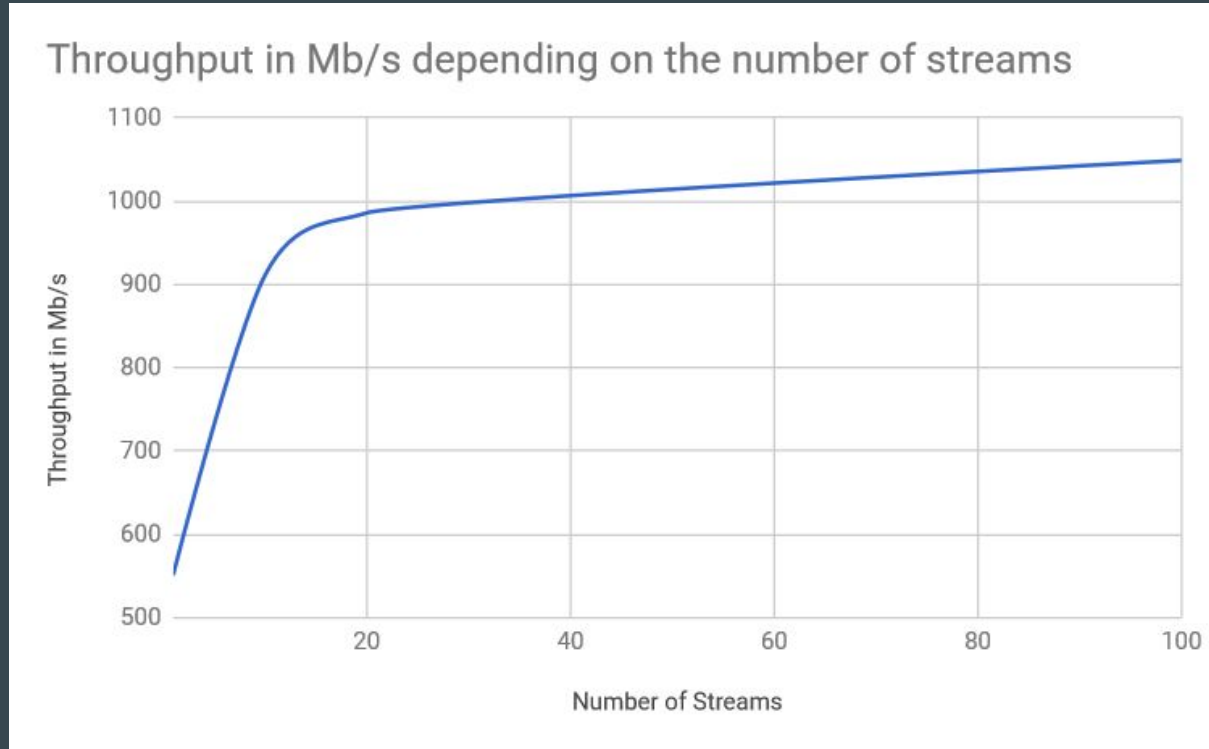
Bandwidth between host and device in Mb/s



Measures



Measures



Measures

BER according to SNR

